

# From reads to results

Dr Torsten Seemann



# What I will cover \*



NGS  
Applications

Sequences  
Sequence quality  
Read file formats  
Using reads  
Alignment file formats  
Analysis tools

RNA-Seq  
DGE



*\* May be different to what you remember tomorrow*

# Reads

# In an ideal world...




- Collect a human genomic DNA sample
- Run it through the lab sequencing machine
- Get back 46 files:
  - phased, haplotype chromosomes
  - each a single contiguous sequence of A,G,T,C
- And it only costs \$1000

# The awful truth



- No such instrument exists
  - can't read long stretches of DNA (yet)
- But we can read short pieces of DNA
  - shred DNA into ~500 bp fragments
  - we can read these reliably
- High-throughput sequencing
  - sequence millions of different fragments in parallel
  - various technologies to do this
  - costs much more than \$1000

# What you get back



## Millions to billions of reads (big files):


```
ATGCTTCTCCGCCTTTAATTAAAATTCATTTTCGTGCACCAACACCCGTTCTTACCATAATAGCTGTTGGAGTCGCTAAACCTAATGCACATGGACACGC
CTAAGATACTGCCATCTTCTTCCAACGTAAATTGTACGTGATTTTCGATCCATTTTCTTCGAGGTTCTACTTTGTCACCCATTAGTGTGGTTACTCGACG
GAATATGCGTGGACAGATGACGAATTGGCAGCAATGATTAAAAAAGTCGGCAAAGGATATATGCTACAGCGATATAAAGGACTTGGAGAGATGAATGCGG
ATCAATGCAAAATACAAGATGTGACAATGCGCGCAATGCAATGATAACTGGTGTGTGCAAAAAGAAACCGAATGTCGTACCTAGTGCAACAGCCACTGCAA
GGAAAAAATGAGAAAAAATTCAGTTCGAAAACTAACGATTTCTGCTTTATTGATTGGGATGGGGTCATTATCCCAATGGTTATGCCTAAAATCATGATC
GATGAAACAATCCAACAAATACCATTCAATAATTTACAGGGGAAAATGAGACNCTAAGTTTCCCCGTATCAGAAGCAACAGAAAAGAAATGGTGTTTCGCT
...
<--- 100 bp --->
...
AGGCATCTTGAAAAACAAGTGTGTGCCTCTGCGATAATCAATGCCACAGAGGTGCATAAAATTAGTTGTCGAAAAATAATCGCTACCGTTGAGACTTC
AAAGGAGCATTCTTCGCACGCGGCAAAAAAGAATACAAACGCATGTCTATAAAGAGACAACCCAAATTACCAGACAGTTAAACGCGATTTATAAGGCT
GTGACAAAAATCGTGTACAGCTTCTTTTATATCCTGTCTTTTTTTAGTTATTTATTTTTCAACCTTATCAATATGACTTGATAGCCTTTTCTTTTCGA
AACTTGTTAAAAAAGACGTCAATGCCTTAACTGTACGTGATTCTTCTGCAGTTAGGGGATGACCTTTGACTACTAAAACAGATGCCATATGCTTACCTTC
ACAAAGCATATTTGTAGGAACGATTGAAAGCATCACTCAAGTAGAAGCGGAAGAAGAAACGATTCAACTGAACTCGTCGATGTCATGGCCAAAGAAGAT
AATTGGACTTTGTACCGATTTTCAGTTTCATCTATGTCCACGCTTATTTTTTCAGCAGTAGCATTCAAAATCACTCCGTCATTGCTGAATGATGTCCCA
CTCCTGTTTCTTTATCTATAATTGAACTGTAAACATGAGGAATCACTTTTTTACACCTGCATCGATTGCAATTTTCAGAATTTCTTCAAAGTTTGAAAG
AAACTGCCATTCAAATGCTGCAAGACATGGGAGGTACTTCAATCAAGTATTTCCCGATGAAAGGCTTAGCACATAGGGAAGAATTTAAAGCAGTTGCGGA
ATCATTCTACGCCAGTCATTTTCGCGTAGTTCTTTTACCATTTTAGCTGTAACGTCTGCCATGTTTAACTCCTCCTGTGTGTGTCTTTTAAAAAAGC
```

<- 1st read

<- 2nd read

<- last read

# Applications



*If you can transform your assay in to sequencing lots of short pieces of DNA, then NGS is applicable.*

## Not just whole genome DNA:

- exome (targeted subsets of genomic DNA)
- RNA-Seq (transcripts via cDNA)
- ChIP-Seq (protein:DNA binding sites)
- HITS-CLIP (protein:RNA binding sites)
- methylation (bisulphite treatment of CpG)
- ... even methods to sequence peptides now!

# FASTA format



# FASTA



```
>NM_006361.5 Homo sapiens homeobox B13 (HOXB13), fragment  
TCTTGCGTCAAGACGGCCGTGCTGAGCGAATGCAGGCGACTTGCGAGCTGGGAGCGAT  
TTGGATTCCCCCGGCCTGGGTGGGGAGAGCGAGCTGGGTGCCCCCTAGATTCCCCGCC  
CCCGGCCGACCCTCGGCTCCATGGAGCCCGGCAATTATGCCACCTTGGATGGAGCCAA  
GGATATCTGGGAGCGGGAGGGGGGCGGAATCTG
```

# FASTA components

Start  
symbol

Sequence ID  
(*no spaces*)

Sequence description  
(*spaces allowed*)



```
>NM_006361.5 Homo sapiens homeobox B13 (HOXB13), fragment
TCTTGCGTCAAGACGGCCGTGCTGAGCGAATGCAGGCGACTTGCGAGCTGGGAGCGAT
TTGGATTCCCCCGGCCTGGGTGGGGAGAGCGAGCTGGGTGCCCCCTAGATTCCCCGCC
CCCGGCCGACCCTCGGCTCCATGGAGCCCGGCAATTATGCCACCTTGGATGGAGCCAA
GGATATCTGGGAGCGGGAGGGGGGCGGAATCTG
```

The diagram illustrates the components of a FASTA sequence entry. It shows a sample entry with four arrows pointing to its parts: a green arrow to the start symbol '>', a red arrow to the sequence ID 'NM\_006361.5', a blue arrow to the sequence description 'Homo sapiens homeobox B13 (HOXB13), fragment', and a brown arrow to the sequence itself, which is split across five lines of 60 characters each.

The sequence  
(*usually 60 letters per line*)

# Multi-FASTA



Concatenation of individual FASTA entries,  
using ">" as an entry separator

```
>read00001
TCTTGCGTCAAGACGGCCGTGCTGAGCGAATGCAGGCGACTTGCGAGCTGGGAGCGA
>read00002
TGGATTCCCCCGGCCTGGGTGGGGAGAGCGAGCTGGGTGCCCCCTAGATTCCCCGCC
>read00003
GGCCGACCCTCGGCTCCATGGAGCCCGGCAATTATGCCACCTTGGATGGAGCCAAGG
>read00004
TCTGGGAGCGGGAGGGGGGCGGAATCTGGAGCGAGCTGGGTGCCCCCTAGATTCCCC
>read00004
GCGGAATCTGGAGCGAGCTGGGTGCCCCCTAGATTCCCCGCATCGTAGATTAGATAT
```



# The DNA alphabet


- Standard
  - A G T C
- Extended
  - adds N (unknown base)
- Full
  - adds R Y M S W K V H D B (ambiguous bases)
  - R = A or G (puRine)
  - Y = C or T (pYrimidine)
  - ... and so on for all the combinations

# Sequence Quality

# Sequences have errors

- nonsense reads
  - *instrument oddness*
- duplicate reads
  - *amplify a low complexity library*
- adaptor read-through
  - *fragment too short*
- indel errors
  - *skipping bases, inserting extra bases*
- uncalled base
  - *couldn't reliably estimate, replace with "N"*
- substitution errors
  - *reading wrong base*

Less common



More common

# Illumina reads



- Usually 100 bp (soon 250 bp)
- Indel errors are rare
- Substitution errors  $< 1\%$ 
  - Error rate higher at 3' end
- Adaptor issues
  - rare in HiSeq (*TruSeq* prep)
  - more common in MiSeq (*Nextera* prep)
- Very high quality overall

# DNA base quality



- DNA sequences often have a *quality value* associated with each nucleotide
- A measure of reliability for each base
  - as it is derived from physical process
    - chromatogram (Sanger sequencing)
    - pH reading (Ion Torrent sequencing)
- Formalised by the *Phred* software for the Human Genome Project



# Phred qualities




Quality	Chance it's wrong	Accuracy	Description
10	1 in 10	90%	Bad
20	1 in 100	99%	Maybe
30	1 in 1000	99.9%	OK
40	1 in 10,000	99.99%	Very good
50	1 in 100,000	99.999%	Excellent

$$Q = -10 \log_{10} P \quad \Leftrightarrow \quad P = 10^{-Q / 10}$$

Q = Phred quality score

P = probability of base call being incorrect

# Quality plot (*FastQC*)



Y-axis is "Phred" quality values (higher is better)

# Quality filtering



- Keep all reads
  - let the downstream software cope
- Reject some reads
  - average quality below some threshold
  - contain any ambiguous bases
- Trim reads
  - remove low quality bases from end
  - keep longest "sub-read" that is acceptable
- Best strategy is analysis dependent

# FASTQ files

# FASTQ



A sequence read looks like this:

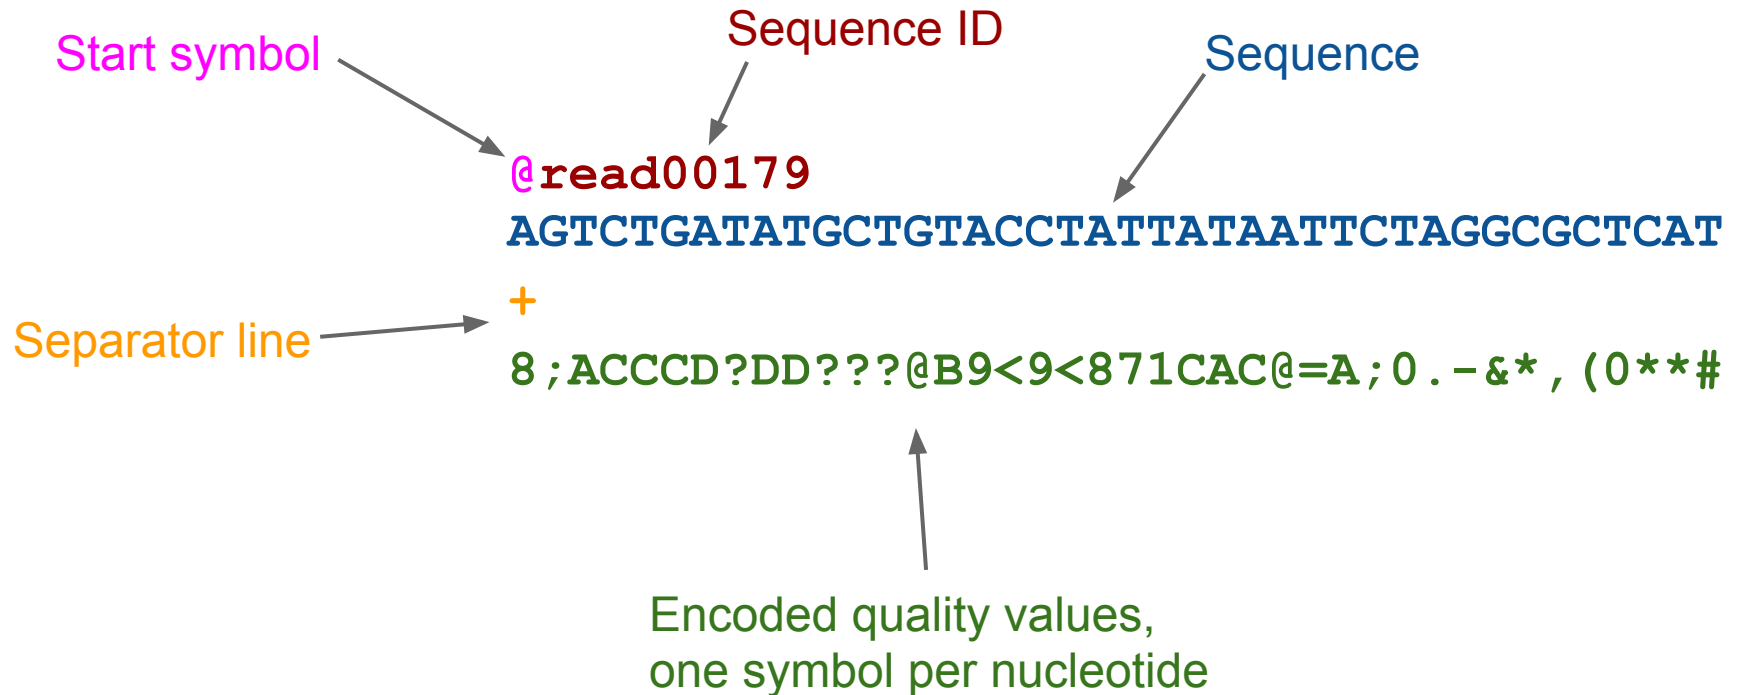
```
@read00179
```

```
AGTCTGATATGCTGTACCTATTATAATTCTAGGCGCTCAT
```


```
+
```

```
8;ACCCD?DD???@B9<9<871CAC@=A;0.-&*,(0**#
```

# FASTQ components



# FASTQ quality encoding



Uses letters/symbols to represent numbers:

! " # \$ % & ' ( ) * + , - . / 0 1 2 3 4 5 6 7 8 9 : ; < = > ? @ A B C D E F G H I J									
Q0		Q10		Q20		Q30		Q40	
<i>bad</i>		<i>maybe</i>		<i>ok</i>		<i>good</i>		<i>excellent</i>	

# Multi-FASTQ



Same as multi-FASTA, just concatenate:

```
@M00267:3:000000000-A0AGE:1:1:15997:1501
CTCGTGCTCTACTTTAGAAAGCTAATGATTCTGTTTGTAGAACATTTTCTACCACTACATCTTTTTCTTGCTTCGCATCTT
+
:=?DD:BDDF>FFHI>E>B9AE>4C<4CCAE+AEG3?EAGEHCGIIIIIIIIIIIGIIIEIIIIIGGIDGIID/;4C<EE
@M00267:3:000000000-A0AGE:1:1:15997:1501
GCCTATAGTAGAAGAAAAAGAAGTGGCTCAAGAAATGAGTGCACCGCAGGAAGTTCCAGCGGCTGAATTACTTCATGAAA
+
<@@FFF?DHFHGHIIIFGIIGIGICDGECHIIIIIIIIIGIHIIFG<DA7=BHHGGIEHDBEBA@CECDD@CC>CCCAC
@M00267:3:000000000-A0AGE:1:1:14073:1508
GTCTTGCTAAATTTAAATAATCTGAAATAATTTGTTCTGCCCGGTCCAATTCAGCTAATACGAGACGCATATAATCCTTA
+
:?DDDDD?84CFHC><F>9EEH>B>+A4+CEH4FFEHFHIIIIIIIIIIIIIGGIIIIIIIIIG>B7BBEBBB@CDDCF
@M00267:3:000000000-A0AGE:1:1:14073:1508
ACGTACAGAGATGCAAAAGTCAGAGAACTTAATATTGTAAGTGAAGTAGCAGCAAGTGTGACATGAGGTTGAAATC
+
1@@DDADHGD?FBGGAFHHCHGGCGGFHIECHGIIGIGFGHGHIIHHEGCCFCB>GEDF=FCFBGGGD@HEHE9=;AD
```



# Data compression



- FASTQ files are very big
  - typically > 10 gigabytes
  - they are somewhat redundant
- Often they will be compressed
  - gzip (.gz extension)
  - bzip2 (.bz2 extension)
  - these are like .ZIP but different method
- Usually get to < 20% of original size
  - faster transfer, less disk space
  - can be slower to read and write though

# FASTQ file name conventions

Suffix	Usage
.fastq .fq	Uncompressed
.fastq.gz .fq.gz	Compressed with GZIP
.fastq.bz2 .fq.bz2	Compressed with BZIP2
s_?_?_sequence.txt	Old Illumina naming (uncompressed)

# Using reads

# In the beginning



First step of most NGS analyses is either:

- ***De novo* assembly**
  - reconstruct the original sequences from reads alone
  - like a jigsaw puzzle but ambiguous
- **Align to reference**
  - find where reads fit on a known sequence
  - can not always be uniquely placed

***De novo* assembly**



# ***De novo assembly***

- *Reconstruct the original DNA sequences using the sequence reads alone*
- Method
  - align each read against every other read
  - build an overlap graph
  - simplify the graph (necessary due to read errors)
  - find distinct paths through the graph
  - calculate the consensus sequence for each path
  - these contiguous sequences are contigs
- Computationally challenging problem

# Why *de novo* ?

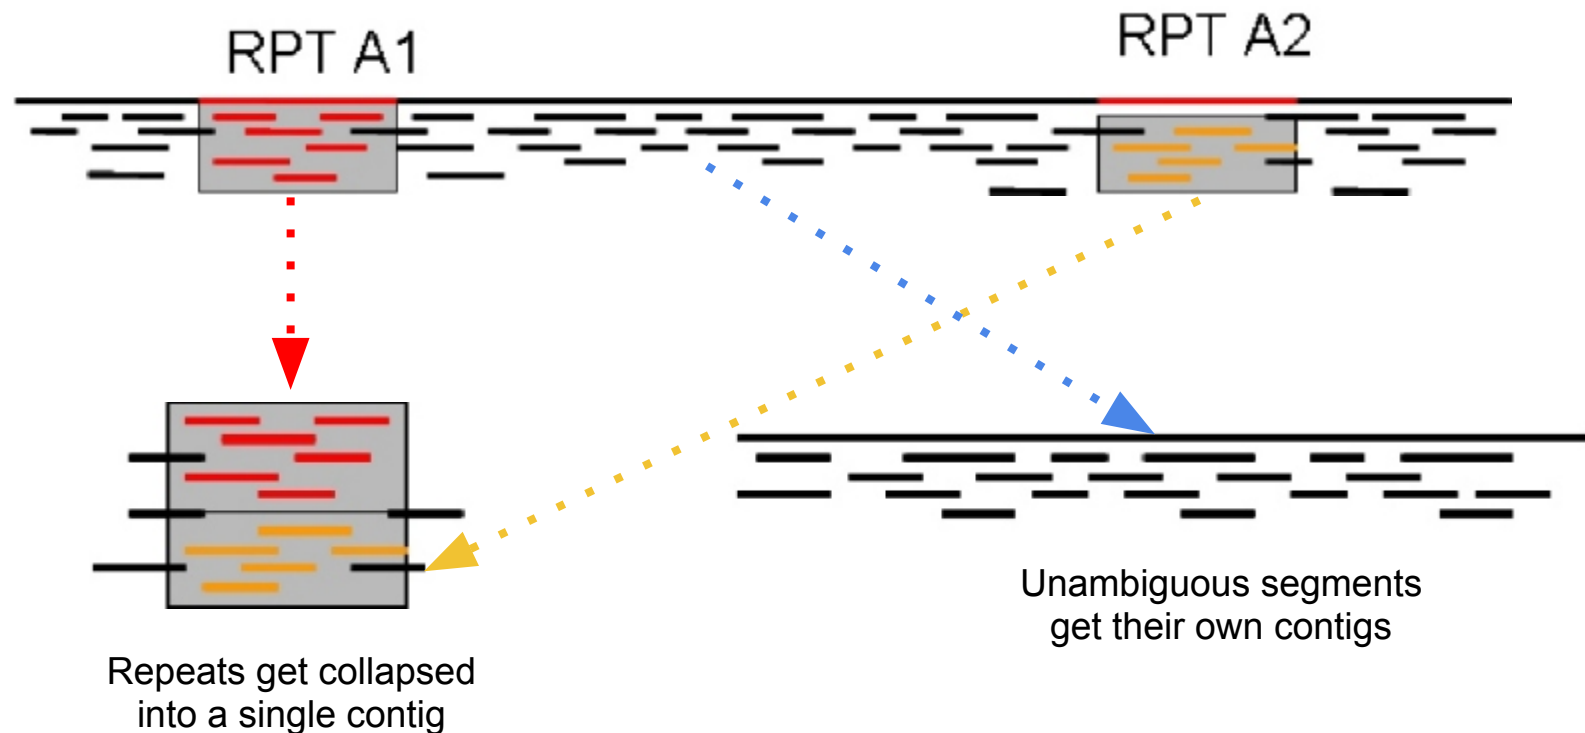


- Sequence a new organism
  - DNA-Seq (genome, expect novel DNA)
  - RNA-Seq (transcriptome, splice variants)
- Unaligned reads from reference alignment
  - novel DNA segments
  - novel RNA transcripts
  - fusion genes
  - contamination

# Where assembly "fails"



*It is impossible to resolve (disambiguate) repeats of length  $L$  with reads shorter than  $L$*






# Assembly file format



- Usually a simple multi-FASTA file of contigs
  - some assemblers provide quality values
  - sequences may contain gaps of "N"s
- Loss of information
  - contigs are a "collapsing" of the rich graph structure
  - ambiguous, but useful connective info is lost
- New standard coming
  - retain all graph information
  - will allow development of new post-processing tools

**Align to reference**

# NGS read alignment



- Want to find where all the reads fit on our reference genome, quickly and accurately.
- Query
  - Lots (>100M) of short (~100bp) reads (FASTQ)
- Reference
  - eg. Human genome, ~27000 contigs (FASTA)
- Many shorts vs. few longs
  - BLAST isn't suitable (it's better at the opposite)
  - New tools: BWA, Bowtie, BFAST, SHRiMP, MAQ

# Example

## Seven short 4bp reads:

AGTC TTAC GGGA CTTT TAGG TTTA ATAG

## The 31bp reference:

AGTCTTTATTATAGGGAGCCATAGCTTTACA

AGTC TAGG ATAG TTAC  
TTTA GGGA CTTT

## Coverage:

1111111100111211100111101122110

Average coverage (depth):  $28 / 31 = 0.90\times$

# Ambiguous alignment



Eight short 4bp reads:

AGTC TTAC GGGA CTTT TAGG TTTA ATAG **TTAT**

The 31bp reference:

**AGTCTTTATTATAGGGAGCCATAGCTTTACA**

AGTC TAGG ATAG TTAC  
TTTA GGGA CTTT  
**TTAT**  
**TTAT** ←

Reads can align to more than one place!

# Multiple-mapping reads



- Align to all possible places
  - useful in some situations
  - but belongs to only one "real" place
- Align to the first place you find
  - not a good idea... but some tools still do it
- Align to a random choice of all valid places
  - useful in some situations
- Don't use multiple-mapping reads
  - often necessary if calling SNPs

# The trade-off



- speed vs. sensitivity
- will miss divergent matches
- can miss indels (insertions and deletions)

# BAM files



# Storing alignments



- SAM
  - plain text file, tab separated columns
  - "a huge spreadsheet"
  - inefficient to read and store
- BAM
  - a compressed version of SAM (~80% less storage)
  - can be indexed (fast access to subsections)
  - needs to be sorted to be useful however
- Standardized format
  - readable by most software

```

1:497:R:-272+13M17D24M    113    1    497    37    37M    15    100338662    0
CGGGTCTGACCTGAGGAGAACTGTGCTCCGCCTTCAG    0;====9;>>>>=>>>>>>>>>>=>>>>>>>>
XT:A:U      NM:i:0      SM:i:37      AM:i:0      X0:i:1      X1:i:0      XM:i:0

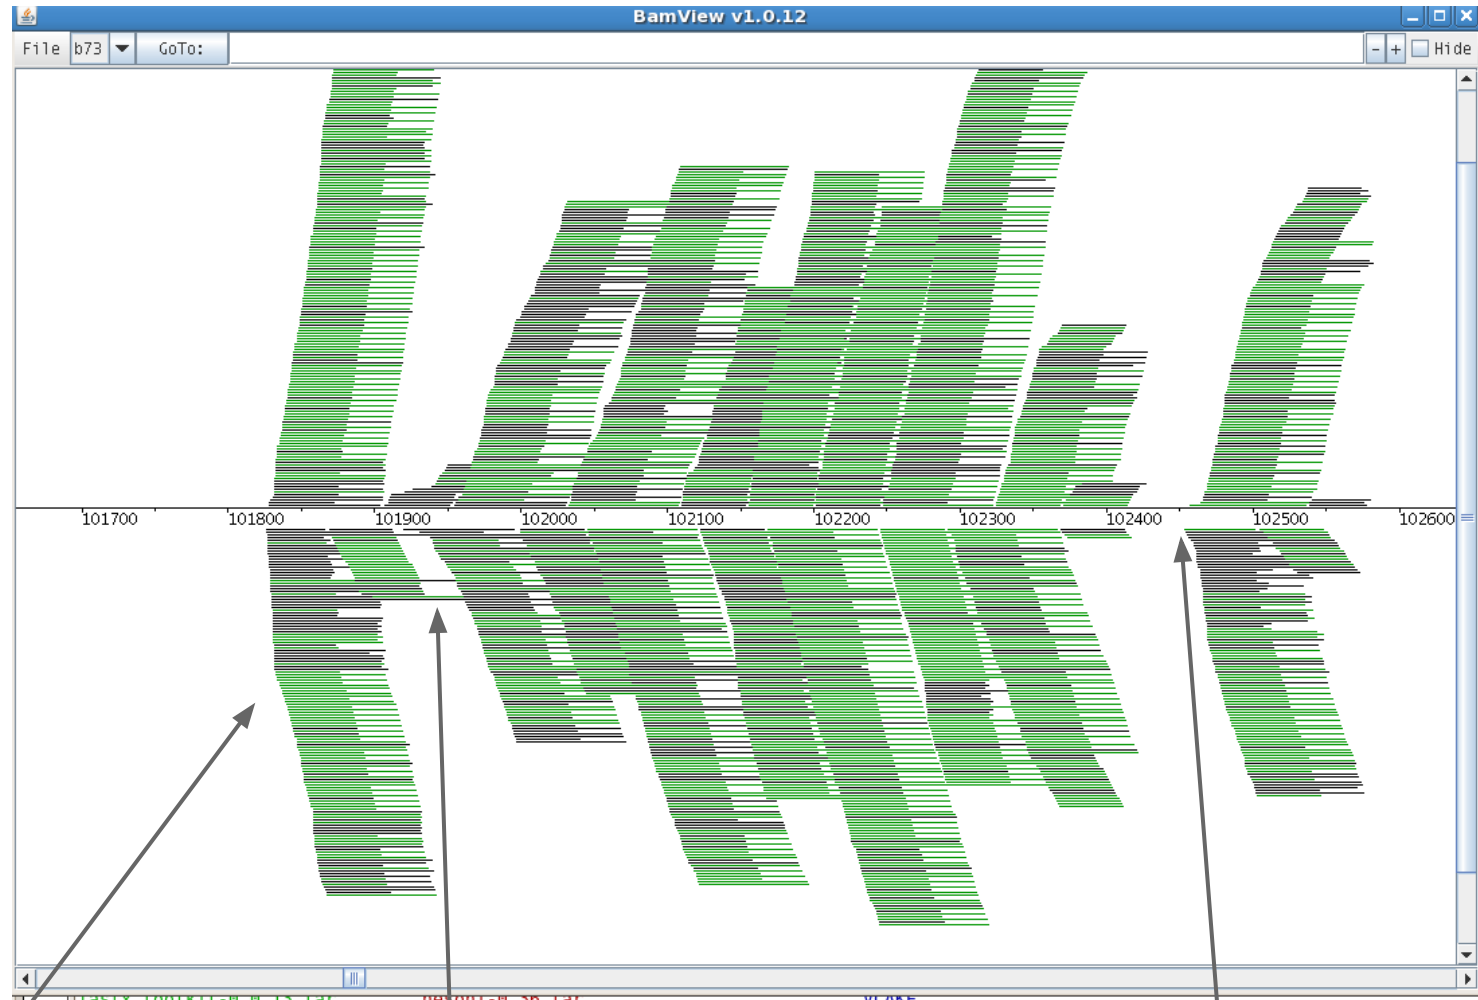
19:20389:F:275+18M2D19M    99    1    17644    0    37M    =    17919    314
TATGACTGCTAATAATACCTACACATGTTAGAACCAT    >>>>>>>>>>>>>>>>>>>><<<<>><<>>4::>>:<9    RG:Z:
UM0098:1    XT:A:R      NM:i:0      SM:i:0      AM:i:0      X0:i:4      X1:i:0      XM:i:0

19:20389:F:275+18M2D19M    147    1    17919    0    18M2D19M    =    17644    -314
GTAGTACCAACTGTAAGTCCTTATCTTCATACTTTGT    ;44999;499<8<8<<<8<<<<<<<<<<7<;<<<>><<    XT:A:R
      NM:i:2      SM:i:0      AM:i:0      X0:i:4      X1:i:0      XM:i:0      MD:Z:18^CA19

```

- One line per original read sequence
  - where it aligned (if at all)
  - how much of it aligned (soft/hard clipping)
  - how well it aligned (mapping quality)
  - any differences to the reference (CIGAR string)
  - lots of other stuff (aligner dependent)

# Wide view (BamView)




High coverage

Low coverage

Zero coverage

# Medium view (IGV)

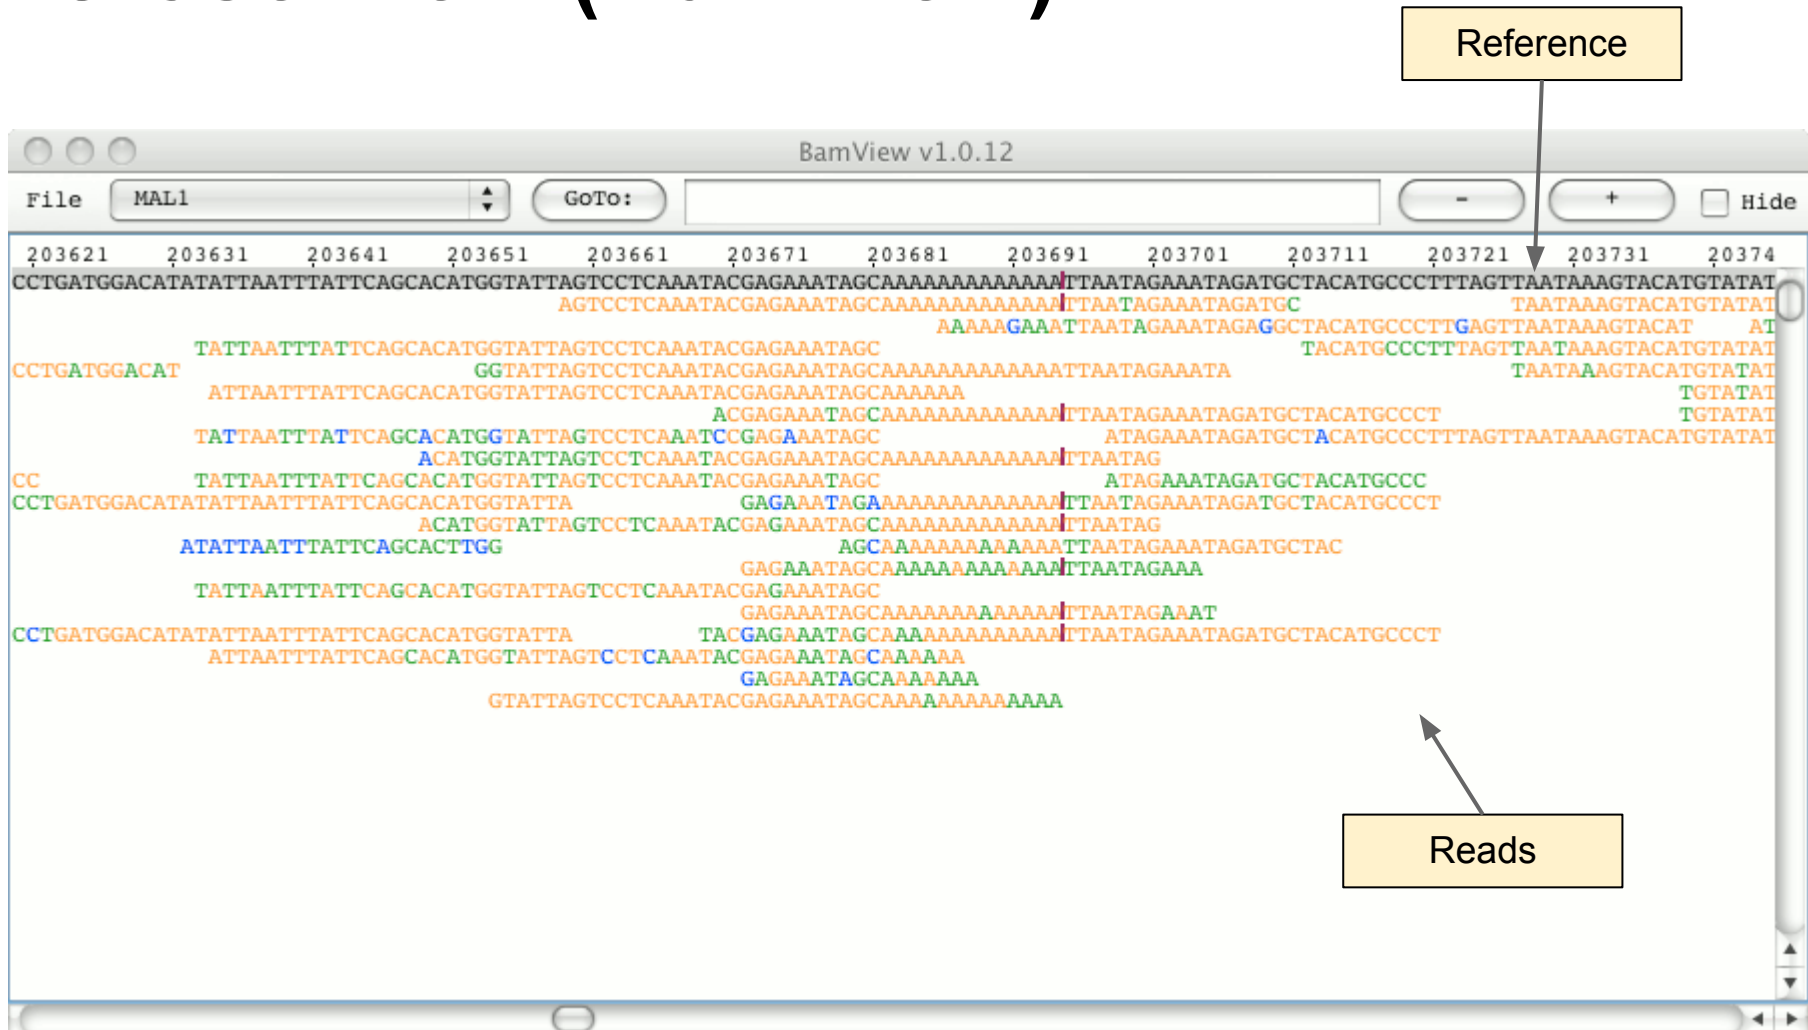


Reads

A variant!

Ref seq

# Close view (BamView)



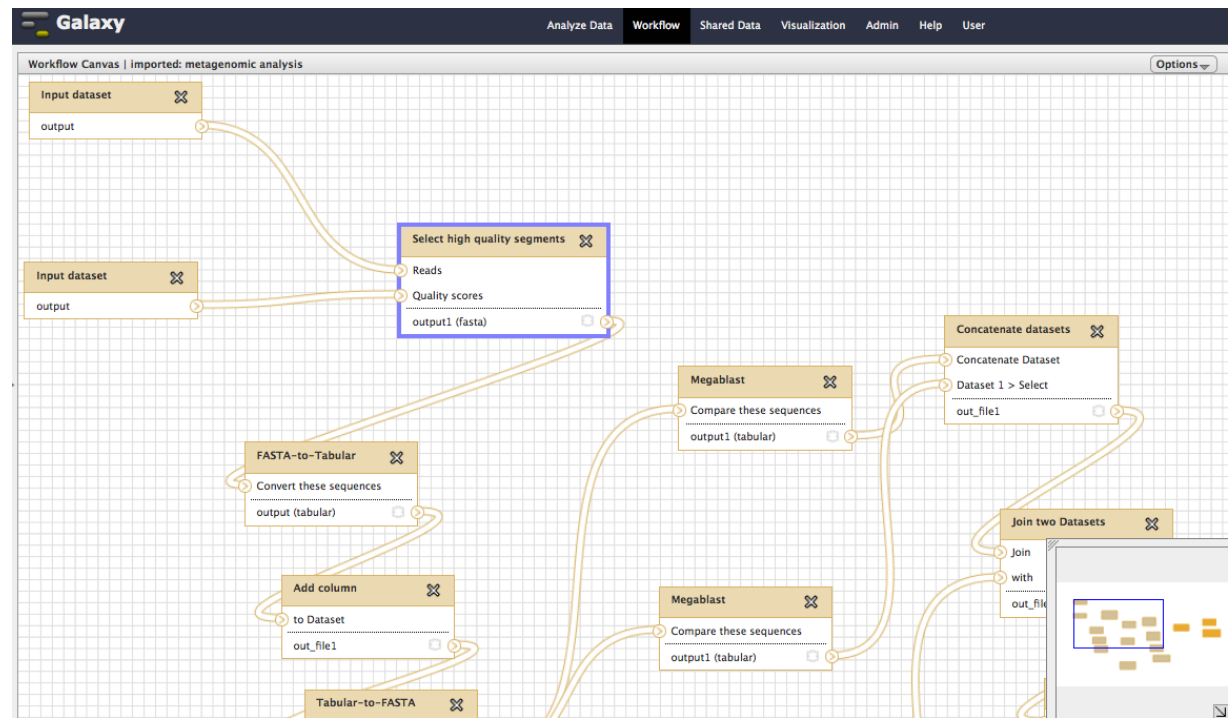
# Getting results

# Analysis



- **Desktop**
  - Free - need Unix skills or a bioinformatician
  - Commercial - CLC Genomics Workbench
- **Server / Cloud**
  - Free - Galaxy / Genomics Virtual Laboratory
  - Commercial - Illumina BaseSpace, ...
- **Service / Subscription**
  - AGRF, QFAB, Geneworks, ....
  - VLSCI::LSCC - hire team at 0.5 EFT etc

# Galaxy



The screenshot shows the Galaxy 101 tutorial page. The header includes the Galaxy logo and navigation tabs: Analyze Data, Workflow, Shared Data, Visualization, Admin, Help, and User. The main content area features the title "Galaxy 101" and the subtitle "Start small The very first tutorial you need". Below the title, there are three "Live Quickies" cards:

- Advanced fastQ manipulation:** Galaxy quickie # 14
- 454 Mapping: Single End:** Galaxy quickie # 15
- Uploading Data using FTP:** Galaxy quickie # 17

The left sidebar contains a "Tools" section with a search bar and a list of tool categories: Get Data, Send Data, ENCODE Tools, Lift-Over, Text Manipulation, Convert Formats, FASTA manipulation, Filter and Sort, Join, Subtract and Group, Extract Features, Fetch Sequences, Fetch Alignments, Get Genomic Scores, Operate on Genomic Intervals, Statistics, Graph/Display Data, Regional Variation, Multiple regression, Multivariate Analysis, Evolution, Motif Tools, Multiple Alignments, Metagenomic analyses, Human Genome Variation, Genome Diversity, EMBOS, NGS TOOLBOX BETA, NGS: QC and manipulation, NGS: Mapping, and NGS: SAM Tools. The right sidebar shows the "History" section with a message: "Your history is empty. Click 'Get Data' on the left pane to start".



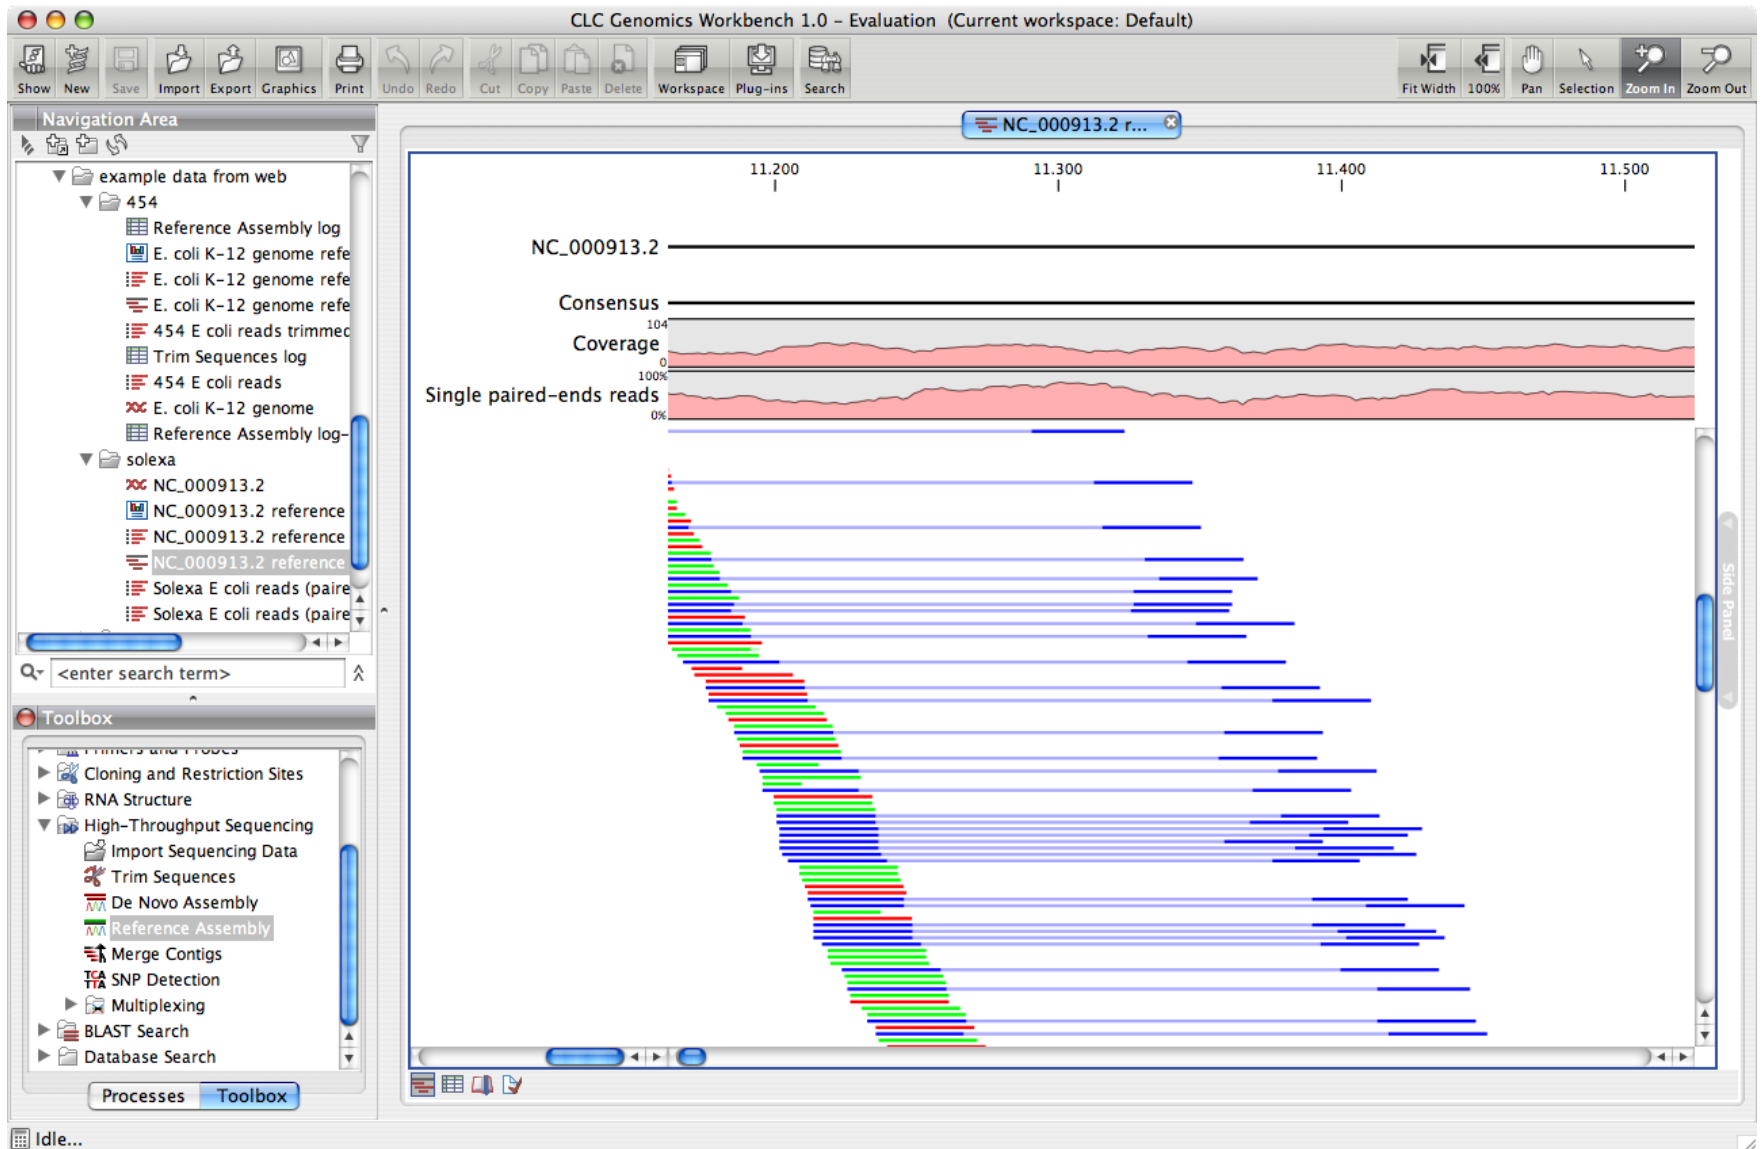


# Galaxy



- Web-based
  - install on your own server
  - install via Cloudman onto your Amazon EC2 cloud
  - use the free public server at *usegalaxy.org*
- Lots of tools and workflows available
  - getting close to plug and play
- Coming soon...
  - Australian Research Cloud
  - will host the Genomics Virtual Laboratory
  - you will be able to run Galaxy instances on it

# CLC Genomics Workbench



# CLC Genomics Workbench



- Designed for NGS analysis
  - Handles most common analyses
    - alignment, RNA-Seq DGE, BLAST, assembly ...
  - but results aren't as good as bioinformatician
    - using Unix and R tools
- Accessible
  - Intuitive interface
  - Runs on Windows, Mac, Linux
  - Needs powerful desktop - lots of RAM
  - "Affordable" licences ~\$4000 (highly variable)

# Summary

# Key points



- Getting into NGS is not trivial
  - new thinking, methods, hardware, software
- Understand the main file types
  - FASTA, FASTQ, SAM, BAM
- Repeat regions cause lots of problems
  - hard/impossible to assemble
  - multi-mapping reads
- No solution fits all problems
  - May need to collaborate/employ bioinformaticians

